

Java Collections Framework

COMS M0103
OOP with Java
Tim Kovacs

2010

Objectives

- To learn what data structures and algorithms the Java library has
- To learn what they can do for you
- To see a few examples of how the classes are meant to be used
 - Sometimes called *idioms*
 - A conventional way of doing something
 - Like a design pattern but at the code level, not the architecture level
- NOT to learn everything you need to know
 - This will let you handle the usual cases
 - For special cases you will need more details

What you should already know

- The concepts of arrays, hash tables and linked lists
- How to use Java's ArrayList and Iterator classes
 - We do not cover basic examples of their use
- A little about the complexity of algorithms
 - that two algorithms with the same functionality can have massively different run-times
 - that run-time efficiency depends a lot on the data
 - So (often) alg. A isn't *always* better than alg. B
 - You need the right tool for the job
- A little about generics (typed collections)
- If you don't know these subjects you can probably still follow this lecture

Contents

- Introduction
- Generic classes
- JCF interface classes*
- JCF implementation classes*
- JCF algorithms*
- Choosing a collection
 - Review of linked lists and hash tables
- Miscellanea
 - Hash codes and equals()
 - New collections syntax in Java 7
 - Collection classes from before the JCF
- Conclusion
 - Summary of key points
 - Where to look for more information

* From the collections trail of the Java tutorial for Java 6

<http://java.sun.com/docs/books/tutorial/collections/index.html>

Collections and the JCF

- A collection is a general-purpose object used to hold and manipulate other objects e.g. ArrayList
- The *Java Collections Framework* (JCF) is Java's collection library
 - Interfaces for different kinds of collections (List, Set...)
 - Implementations of the interfaces
 - Algorithms for manipulating collections (searching, sorting...)
- Implements efficient collections and algorithms so you don't have to
- The JCF is complex and we don't cover everything!
- Don't confuse with the JFC (Java Foundation Classes)
 - A set of GUI packages: AWT, Swing and Java 2D

Some JCF Design Patterns

- Views
 - Viewing (part of) a collection using another collection
 - Possibly a different collection type
 - E.g. getting a sub-list of a List
 - E.g. getting the set of keys in a Map as a Set
 - The view shares its contents with the original
 - More efficient than copying the contents
 - Lets you manipulate contents of the original
- Factory pattern (for iterators)
 - Getting a collection to construct its own iterator
 - Means you don't need to know the correct type of iterator
- Optional interface methods
 - A messy way of not implementing all the functionality specified by an interface
- Strategy pattern
 - Comparator objects used by sorting algorithms

Generic Classes

Generic Classes

- The JCF uses a lot of generic classes
 - Classes with a *type parameter*
 - E.g. `ArrayList<String>` instead of just `ArrayList`
 - The two are different types!

```
ArrayList<String> myList = new ArrayList(); // error
```

```
ArrayList<String> myList = new ArrayList<String>(); // ok
```

- A type parameter is passed to a class
 - Just like an ordinary parameter is passed to a method
- Why not just use type `Object`?
 - Using more specific types lets the compiler check for errors
- See basic generics tutorial

<http://java.sun.com/docs/books/tutorial/java/generics/index.html>

- There's a much more advanced version on same site ⁸

A Generic Class

```
public class Box<T> {           // T stands for "Type"
    private T data;             // define a field of type T, with name "data"

    public void add(T data) {
        this.data = data;
    }

    public T get() {
        return data;
    }
}
```

Now we can create boxes which hold different types

```
Box<Integer> integerBox = new Box<Integer>(); // T = Integer
Box<String> stringBox = new Box<String>();    // T = String
```

Type Parameter Names

- When writing a generic class (or documentation for it) we give the type parameter a name
- Typical names and what they stand for:
 - E - Element in a Collection (used extensively by the JCF)
 - K – Key in a Map
 - N – Number
 - T – Type
 - V – Value in a Map
 - S,U,V etc. - 2nd, 3rd, 4th types
- These are just conventions
 - Like ordinary variables, the name could be (almost) anything
 - Note the convention of 1 letter each
 - Unlike ordinary variables where 1 letter names are considered bad (except on slides!)

More Complex Type Parameters

- Sometimes we have more than one type parameter
 - To put a key/value pair in a map:

```
V put (K key, V value)
```

↑
return type of method put is V

- The wildcard ? means 'any type'
- Sometimes we want to restrict types

```
addAll(Collection<? extends E> c)
```

 - ? can be any subtype of E
 - We call this a *bounded type parameter*
 - We can use `super` instead of `extends`
- There are a lot more details to generics
 - e.g. they only exist at compile-time!
 - At run-time `List<String>` is just `List`
 - See generics tutorials for more


JCF Interfaces

Interfaces as ADTs

- The interface of a module is its set of public methods
- We have seen the utility of hiding objects behind interfaces
 - Polymorphic algorithms: write code using interface type, choose from different implementing types
- The interface for a data structure is sometimes called an Abstract Data Type (ADT)
 - Same principle: write polymorphic code, choose implementing data structure later
 - Choose the right tool (data structure) for the job (data)

Core Collection Interfaces

The most important interfaces in the JCF



- All are generic i.e. collections are typed
 - `Collection<String>`, `Collection<Integer>` ...
 - Since documentation applies to all types we use a variable for the type `Collection<E>`
- We will not look at the Queue interface

Core Collection Interfaces

- Collection
 - JCF has no implementation; only its children implemented
 - Defines core methods for its children
- Set
 - No duplicate elements: `assert !e1.equals(e2)`
 - Elements unordered
- List
 - Can have duplicates. Elements ordered
- Queue
 - For temporary storage and sorting of dynamically arriving objects
 - E.g. First-In, First-Out (FIFO) protocol
- Map
 - Maps keys to values (like a phone book)
 - A key maps to either 0 or 1 value
 - Cannot have duplicate keys

Typical Uses

- Collection
 - Most cases where you just need to store or iterate through a "bunch of things" and later iterate through them
 - Use List if you care about their order
- Set
 - Remembering which items you've already processed
 - e.g. when doing a web crawl
 - Making other yes-no decisions about an item, e.g.
 - is the item an English word?
 - is the item in the database?
 - is the item in this category?

Based on http://www.javamex.com/tutorials/collections/how_to_choose.shtml

Typical Uses

- Map
 - Used in cases where you need to say "for a given X, what is the Y"?
 - For a given user ID, what is the name?
 - For a given IP address, what is the country code?
 - For a given string, how many instances have I seen?
- Queue
 - For traversing hierarchical structures such as a filing system, or in general where you need to remember "what data to process next", whilst also adding to that list of data

Optional Methods

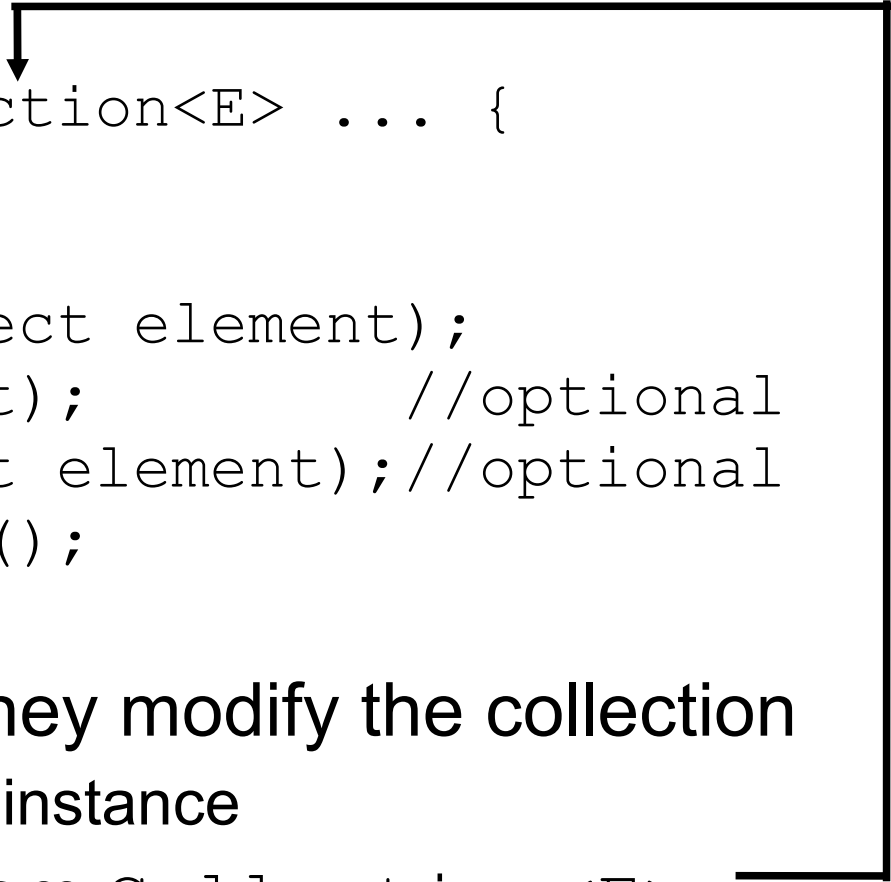
- Some collection methods are marked “optional” in the documentation
 - But the general-purpose implementations we will see later all implement all the optional methods
- Other implementations may throw an `UnsupportedOperationException` instead
 - This is because a class must provide code for all methods in an interface; it can't choose a subset
 - Throwing this exception is a hack to allow them to provide functionality for only a subset

Optional Methods

- In general, this is a **bad** design idea
 - The alternative was to have many more interfaces
 - Versions with and without the optional methods
 - The JCF is already very complex so optional methods seemed a better compromise

The Collection Interface

- Very abstract and general since it's the parent of most of the others
- Basic operations:



```
public interface Collection<E> ... {  
    int size();  
    boolean isEmpty();  
    boolean contains(Object element);  
    boolean add(E element);           //optional  
    boolean remove(Object element); //optional  
    Iterator<E> iterator();  
    ... }
```

- add, remove **return true** if they modify the collection
 - remove removes only a single instance
- The E parameter to add is from Collection<E>
 - In other words you can only add elements of type E

Identity and Equality

- Recall identity and equality:
 - Identity: 2 objects are `==` if they are really the same object
 - I.e. two variables refer to the same object
 - Equality: 2 objects are `equal()` if their state is equivalent

```
Rectangle r1 = new Rectangle(0,0,5,5);
Rectangle r2 = new Rectangle(0,0,5,5);
r1.equals(r2)    // true – same position and size
r1 == r2        // false – different objects
```
- **Class Object defines `equals()` using `==`**
 - So by default `==` and `equals()` are the same
- If we want different rectangle objects to have equality (as in code above) we must override `equals()`

Equality in the JCF

- The contract of many methods uses `equals()`
 - E.g. `contains(element)` returns true if the collection has some object for which `equals(element)` returns true
- This is the method's contract, not a description of implementation
 - A shortcut may be used instead of calling `equals` on all elements

When to use the Collection Interface

- Use type `Collection` when you don't care what the subtype is
 - E.g. when you just want to pass collections around
- E.g. all general-purpose collection implementations have a constructor which accepts type `Collection`
- Called a *conversion constructor* because it converts any collection into its own type
- E.g. suppose you have objects of type `List<String>` and `Set<String>` you want to convert to `ArrayList<String>`
- `ArrayList` only needs one constructor to handle both

```
...new ArrayList<String>(myList);  
...new ArrayList<String>(mySet);
```

Traversing Collections

- Two ways:
 - For-each loops
 - Iterator objects

- For-each loop example

```
for (Object current : myCollection)
    System.out.println(current);
```

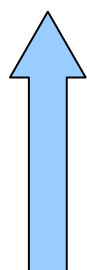
- Limitations
 - You cannot add or remove elements during traversal
 - You cannot alternate iteration over multiple collections

Iterators

```
public interface Iterator<E> {  
    boolean hasNext();  
    E next();  
    void remove(); //optional  
}
```

- **Note: no add method**
 - Unlike ListIterator interface
- **Example**

```
Iterator it = myCollection.iterator();  
while (it.hasNext())  
    if (noLongerNeeded(it.next()))  
        it.remove();
```



Collection's iterator method constructs an iterator for us. An example of the factory pattern.

Traversing Collections

- Use Iterators when you want
 - to remove elements during traversal, or
 - to alternate iteration over multiple collections
- Remember:
 - `remove()` removes the last element returned by `next()`
 - `remove()` can only be used once for each element
 - Using Collection's `add()` or `remove()` breaks existing iterators
 - Use `Iterator.remove()` instead
 - Using `Iterator.remove()` breaks any other iterators.
- Rules for using iterators
 - If you don't add or remove you can have multiple iterators
 - If you do remove you should have only 1 iterator

Iterators and the Factory Pattern

- Why not call the iterator constructor yourself?
 - To keep the code generic/abstract/polymorphic
 - Each subclass of Collection has its own iterator class
 - You do not want to tie the code to a particular collection type
 - Even if that didn't matter
 - You probably don't know the name of the right iterator class
 - It's hard to find out
 - E.g. what's the iterator class for a HashSet?

```
Collection myCollection = new HashSet();  
??? it = new ???(); //HashSetIterator?  
// keep it generic; let Java figure it out  
Iterator it = myCollection.iterator();
```

The *only* place we specify what type of collection it is

Bulk Operations on Collections

- Bulk operations (except clear) take a collection as parameter

```
boolean containsAll(Collection<?> c);  
boolean addAll(Collection<? extends E> c); //optional  
boolean removeAll(Collection<?> c); //optional  
boolean retainAll(Collection<?> c); //optional  
void clear(); //optional
```

- booleans indicate whether collection was modified
- Note bounded type parameter to addAll

Idiom: Removing All Copies of 1 Object

- `remove(Object o)` only removes *one* instance of `o`
- `removeAll(Collection c)` removes *all* instances of all elements of `c`
- What if you want to remove all instances of one object `o`?
 - First put `o` in a collection
 - Then use `removeAll`:

```
c.removeAll(Collections.singleton(e));
```
 - `singleton()` is factory method returning an immutable Set with one element
 - Efficient because the Set is a special lightweight class, not a regular Set class

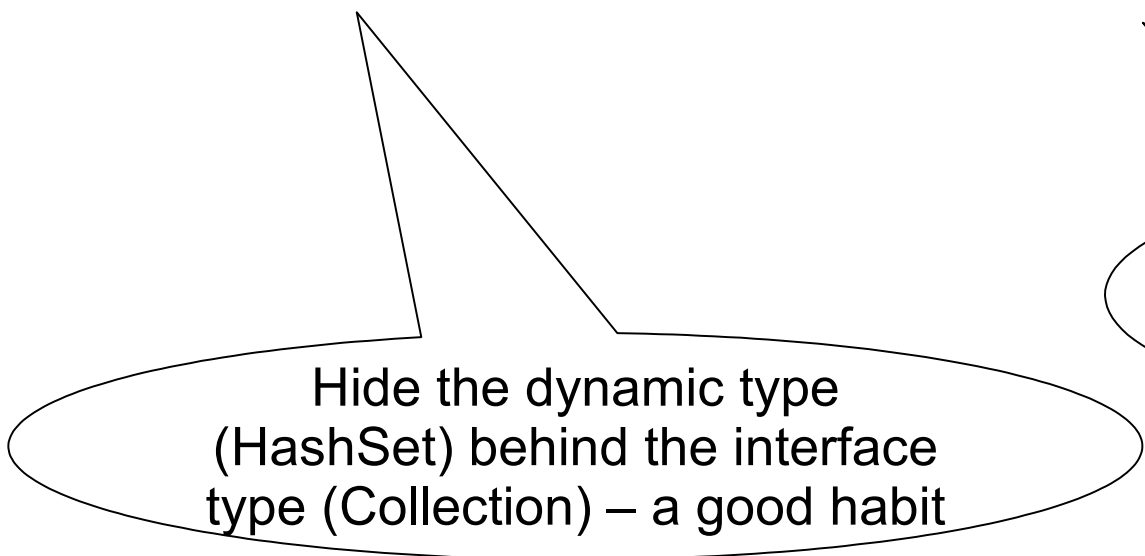
The Set Interface

- Only has methods inherited from Collection
- Cannot contain duplicate elements
 - Can only contain one `null` value
 - Object duplicates are defined as: `e1.equals(e2)`
 - If you try to add a duplicate it ignores it
- `Set A equals()` set B if they contain exactly the same elements
- If you change an element so it becomes a duplicate, behaviour of the set is undefined
 - In other words: don't do it!
 - So be careful with sets that contain mutable elements
- A set cannot contain itself

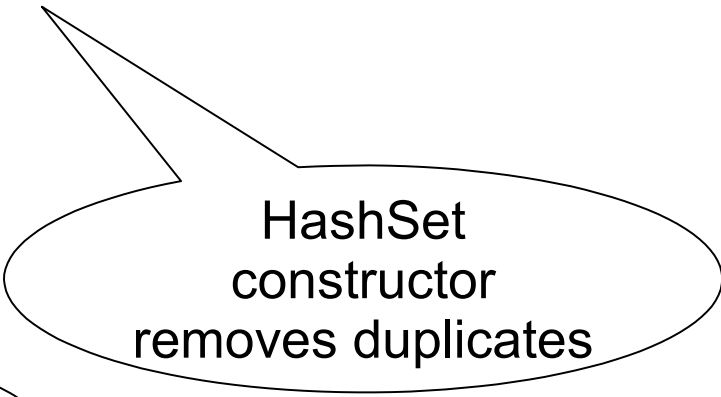
Idiom for Removing Duplicates

- Given collection of strings c1 make a collection c2 which is the same as c1 but without duplicates

```
Collection<String> c2 = new HashSet<String>(c1);
```



Hide the dynamic type
(HashSet) behind the interface
type (Collection) – a good habit



HashSet
constructor
removes duplicates

The List Interface

- Lists are ordered and can contain duplicates
 - Like arrays the first element is at position 0
 - Lists can contain themselves but this can cause problems!
- List inherits from Collection and adds:
 - Random access (also called positional access)
 - Methods like get, set, add, remove which take an index
 - Search methods
 - indexOf and lastIndexOf
 - ListIterator
 - an extended version of Iterator with forward and backward traversal
 - also has an add method so you can add at a position
 - subList(int from, int to)
 - returns a new list which shares objects with the original

Notes on List

- `add` and `addAll` add to the end of the list
 - An iterator can add at other positions
- `remove` removes the first instance
 - Collection's version removes a arbitrary instance
- The only way to get an iterator which points at a particular object is to traverse the list
 - `contains` returns true or false, not a position

Implementations of the List Interface

- The JFC has 2 implementations of List
 - ArrayList and LinkedList
- They have very different run-time behaviour
 - ArrayList has random access
 - LinkedList has sequential access
- Making List fit both of them means
 - List is not very cohesive
 - For each implementation, there are methods in List which are very inefficient
- More in section on implementations

The Map Interface

- Maps keys to values
 - Each key maps to 0 or 1 values
 - Cannot contain duplicate keys
 - Maps are not ordered (but SortedMaps are)
 - “Put” with an existing key replaces the key's last value (and returns it, or `null` if there was no existing value)

```
public interface Map<K,V> {  
    V put(K key, V value);  
    V get(Object key);  
    V remove(Object key);  
    boolean containsKey(Object key);  
    boolean containsValue(Object value);  
    int size();  
    boolean isEmpty();  
    ...}
```

Map's Collection Views

- Interaction with a Map is limited
 - Basically: put(), get(), containsKey(), containsValue()
 - No (direct) way to iterate through all keys
- Collection view methods return part of a map as another Collection type
 - View shares the contained objects with the map
 - Allows viewing the map's contents
 - But also processing them (not read-only)
 - The only way to iterate through a map

Map's Collection Views

```
public interface Map<K,V> {...  
    Set<K> keySet();  
    Collection<V> values();  
    Set<Map.Entry<K,V>> entrySet();  
    ...}
```

- Each method returns a collection:
 - `keySet` — the `Set` of keys in the map
 - `values` — The `Collection` of values in the map
 - Not a `Set` as multiple keys can map to a value
 - `entrySet` — the `Set` of key/value pairs in the map
 - The pairs have type `Map.Entry`
 - `Map.Entry` is an inner interface of `map`
- All are public

Iterating through Keys or Values

Printing all keys with a for-each loop

```
for (<Type> key : myMap.keySet())  
    System.out.println(key);
```

Using an iterator to filter out all elements with some property

```
Iterator<Type> it = myMap.keySet().iterator();  
while (it.hasNext())  
    if (it.next().isBogus())  
        it.remove();
```

Map.Entry

- An interface defined inside Map
- Key/value pairs are stored with this type

```
public interface Entry {  
    K getKey();  
    V getValue();  
    V setValue(V value);  
}
```

Iterating through Key/Value Pairs

```
// print all key/value pairs
for (Map.Entry<KeyType, ValType> e : m.entrySet())
    System.out.println(e.getKey() + ": " +
        e.getValue());
```


Removing, Setting but NOT Adding

- If you remove from a collection view
 - you also remove from its map (since they share their contents)
 - you break any existing iterators of both map and view
 - Except the one that did the removal, if you did it with an iterator
 - See example a few slides ago of filtering out elements
- If you remove a key, you remove its value too
- You can use a view's `Entry.setValue()` to change the underlying map
- You CANNOT add elements to the map using collection views
 - Just add them directly to the map

Multimaps

- Like a map but each key has multiple values
- Not implemented in the JCF
- Just use a Map and for each value use another collection e.g. a List

Ordering Objects

- A set of objects can be ordered (sorted) in many ways
 - a,b,c
 - b,a,c
 - and so on...
- In a *total ordering* it is possible to compare any pair of elements
 - E.g. the set of integers; any two integers can be compared
 - Elements can be tied, e.g. A and a
- In a *partial ordering* some pairs cannot be compared
 - E.g. ordering letters and other symbols alphabetically
 - How do you compare a and @?
 - Of course you can define an ordering (or many) to do this
 - But that goes beyond alphabetical ordering

Comparable Interface

```
public interface Comparable<Type> {  
    public int compareTo(Type other);  
}
```

- Used to sort objects
- Compares objects “this” and “other” and returns:
 - negative integer if this < other
 - 0 if they are equal
 - positive integer if this > other
- Comparable's ordering is called the class's *natural order*
- Implemented by all number types, String, Date and some other library classes

Comparator Interface

- What if you want:
 - to choose different orderings?
 - E.g. sort Person objects by name or date of birth?
 - to sort objects that don't implement Comparable?
- Write a class which implements the **Comparator** interface (not **Comparable**)

```
public interface Comparator<Type> {  
    int compare(Type obj1, Type obj2);  
}
```
- Pass an instance of this class to other classes e.g. SortedSet
 - An example of the *strategy* pattern

“Consistent with Equals”

- Comparable & Comparator can only return 0 if the elements are `equal()`
 - We say the comparison must be “consistent with equals”

SortedSet Interface

- A set that keeps its elements in ascending order
 - Uses natural order, or a Comparator passed to constructor
 - Must be a total ordering
 - Furthermore, since it's a set it cannot have duplicates
 - So if compare or compareTo return 0 the second object will not be inserted
- Extends Set interface to add
 - Range-views
 - Return another SortedSet sharing subset of elements with the first
 - The subset is a consecutive range
 - Get methods for the endpoints (first and last element)
 - A get method for the Comparator it uses

SortedSet Interface

```
public interface SortedSet<E> extends Set<E> {  
    // Range-views  
    SortedSet<E> subSet(E fromElement, E toElement);  
    SortedSet<E> headSet(E toElement);  
    SortedSet<E> tailSet(E fromElement);  
  
    // Endpoints  
    E first();  
    E last();  
  
    // Comparator access  
    Comparator<? super E> comparator();  
}
```


SortedMap Interface

- A map that keeps its elements in ascending order
 - Uses natural order, or a Comparator passed to constructor
 - Sorting is done only on keys, not values
 - Must be a total ordering
- SortedMap is to Map exactly what SortedSet is to Set
 - I.e. extends Map interface to add:
 - Range-views
 - Get methods for the endpoints (first and last element)
 - A get method for the Comparator it uses

```
public interface SortedMap<K, V> extends Map<K, V>{  
    SortedMap<K, V> subMap(K fromKey, K toKey);  
    SortedMap<K, V> headMap(K toKey);  
    SortedMap<K, V> tailMap(K fromKey);  
    K firstKey();  
    K lastKey();  
    Comparator<? super K> comparator();  
}
```

Adding to Sorted Collections

- SortedSet and SortedMap impose extra conditions on the elements they accept
 - Elements must either:
 - implement Comparable, or
 - be accepted by the Comparator object provided
 - Furthermore, all objects must be *mutually comparable*
 - If a collection contains only Rectangles, it only needs to compare a Rectangle and another Rectangle
 - But if it contains Rectangles and Triangles, it must also know how to compare a Rectangle and Triangle
 - You may need to extend Comparable or Comparator for this

JCF Implementations

About Implementations

- We have seen the JCF interfaces
- Each can be implemented by multiple JFC classes
 - or by user-supplied classes
- Implementations can have very different run-time characteristics
 - They may also sort elements in different ways

Kinds of JCF Implementations

- General-purpose implementations
 - For everyday use
- Special-purpose imp.
 - for special purposes (!)
- Concurrent imp.
 - for use with threads
- Wrapper imp.
 - to add or remove functionality from others (Decorator pattern)
- Convenience imp.
 - simple implementations for special purposes
- Abstract imp.
 - partial implementations to help write your own

General-purpose Implementations

All general-purpose implementations:

- implement all optional methods
- allow null elements, keys and values
- are unsynchronised (i.e. not suitable for threads)
- are serializable (can be written to files)
- have public clone methods to make copies
- have *fail-fast* iterators
 - If the collection is modified directly (not through an iterator) the iterator throws `ConcurrentModificationException`
 - But only if iterator detects the modification
 - Better than unpredictable behaviour of iterating over modified collection (“fail later”)

General-purpose Implementations

All General-purpose Implementations

Interfaces	Implementations					
	Hash table	Resizable array	Tree	Linked list	Hash table + Linked list	Priority Heap
Set	HashSet		TreeSet		LinkedHashSet	
List		ArrayList		LinkedList		
Map	HashMap		TreeMap		LinkedHashMap	
Queue				LinkedList		PriorityQueue

- Names are in 2 parts: implementation/interface

Selected General-purpose Implementations

Selected General-purpose Implementations

Interfaces	Implementations			
	Hash table	Resizable array	Tree	Linked list
Set	HashSet		TreeSet	
List		ArrayList		LinkedList
Map	HashMap		TreeMap	

- These are the ones you should know

JCF Set Implementations

- HashSet
 - No guarantees about order of traversal
 - Most efficient version (uses a hash table)
 - Default initial size is 16
 - Rule of thumb: initialise to twice size you expect to need
- TreeSet
 - Visits elements in order of their value
 - Much slower than HashSet (uses a red-black tree)
- LinkedHashSet
 - Visits elements in order of their insertion
 - Slightly slower than HashSet (uses a hash table and a linked list)

JCF Map Implementations

- HashMap, TreeMap and LinkedHashMap
- Just like the corresponding Set implementations

JCF List Implementations

- The JFC has 2 implementations of List
 - ArrayList and LinkedList
- They have very different run-time behaviour
 - ArrayList has random access
 - LinkedList has sequential access
- Making List fit both of them means
 - List is not very cohesive
 - For both ArrayList and LinkedList there are methods in List which are very inefficient
 - Of course the smaller the collection the less it matters
 - Inefficient use is probably only noticeable if you have hundreds of elements and process them intensely
 - But if the collection is big it can make a huge difference

Random vs. Sequential Access

- Random access
 - Jumps straight to any element (think of an array)
- Sequential access
 - Starts at one end and visits every element until it reaches the one you want (think of a linked list)
 - If you want element 1000 you must visit the 1000 elements that come before it – slow!
- List methods with an index argument
 - are fast with ArrayList
 - are very slow with LinkedList

How ArrayList is Implemented

- ArrayList delegates storage to a private array
 - This is called the *backing* array
 - Because it supports (backs up) the ArrayList
- You may not have known that, even if you've used it
 - It's hidden behind the public interface of ArrayList
- Usually it's good to separate the interface and implementation
 - But sometimes it is useful to know about the implementation
 - E.g. so you can choose one that is efficient for the current problem

Efficiency of Insertion and Deletion

- If the array is not full, adding to the end is efficient
- But insertions and deletions which cause reordering are inefficient
 - If the array has 1000 elements and you insert at position 500 it has to shift all elements between 500 to 1000
- If the array is too small
 - A bigger one is created
 - The smaller one is copied into it
- In contrast, insertion and deletion in a LinkedList always involve 1 element
 - Very efficient

LinkedList or ArrayList?

- Use the one that suits your problem
 - Usually ArrayList
- If you're not sure try them both
 - Easy if your code treats them both as type List
- The JCF could have split List into two interfaces
 - One for linked lists and one arrays (plus a List parent)
 - Then it would be clearer how to use them efficiently
 - But they wanted to minimise the number of interfaces

“If you think you want to use a LinkedList, measure the performance of your application with both LinkedList and ArrayList before making your choice; ArrayList is usually faster.”

Josh Bloch, primary JCF designer
Quote from the Collections Trail

Other JCF Implementations

- We have seen only general-purpose ones so far
 - Here are a few others
 - They may not implement the “optional” methods

- EnumSet

- High-performance Set for Enums
 - Allows easy iteration over ranges

```
for (Day d : EnumSet.range(Day.MONDAY, Day.FRIDAY))  
    System.out.println(d);
```

- EnumMap

- High-performance Map using Enums as keys

- Synchronization Wrappers

- Add synchronization (thread safety) to Collection classes

- Unmodifiable Wrappers

- Provide read-only access to Collection classes

JCF Algorithms

JCF Algorithms

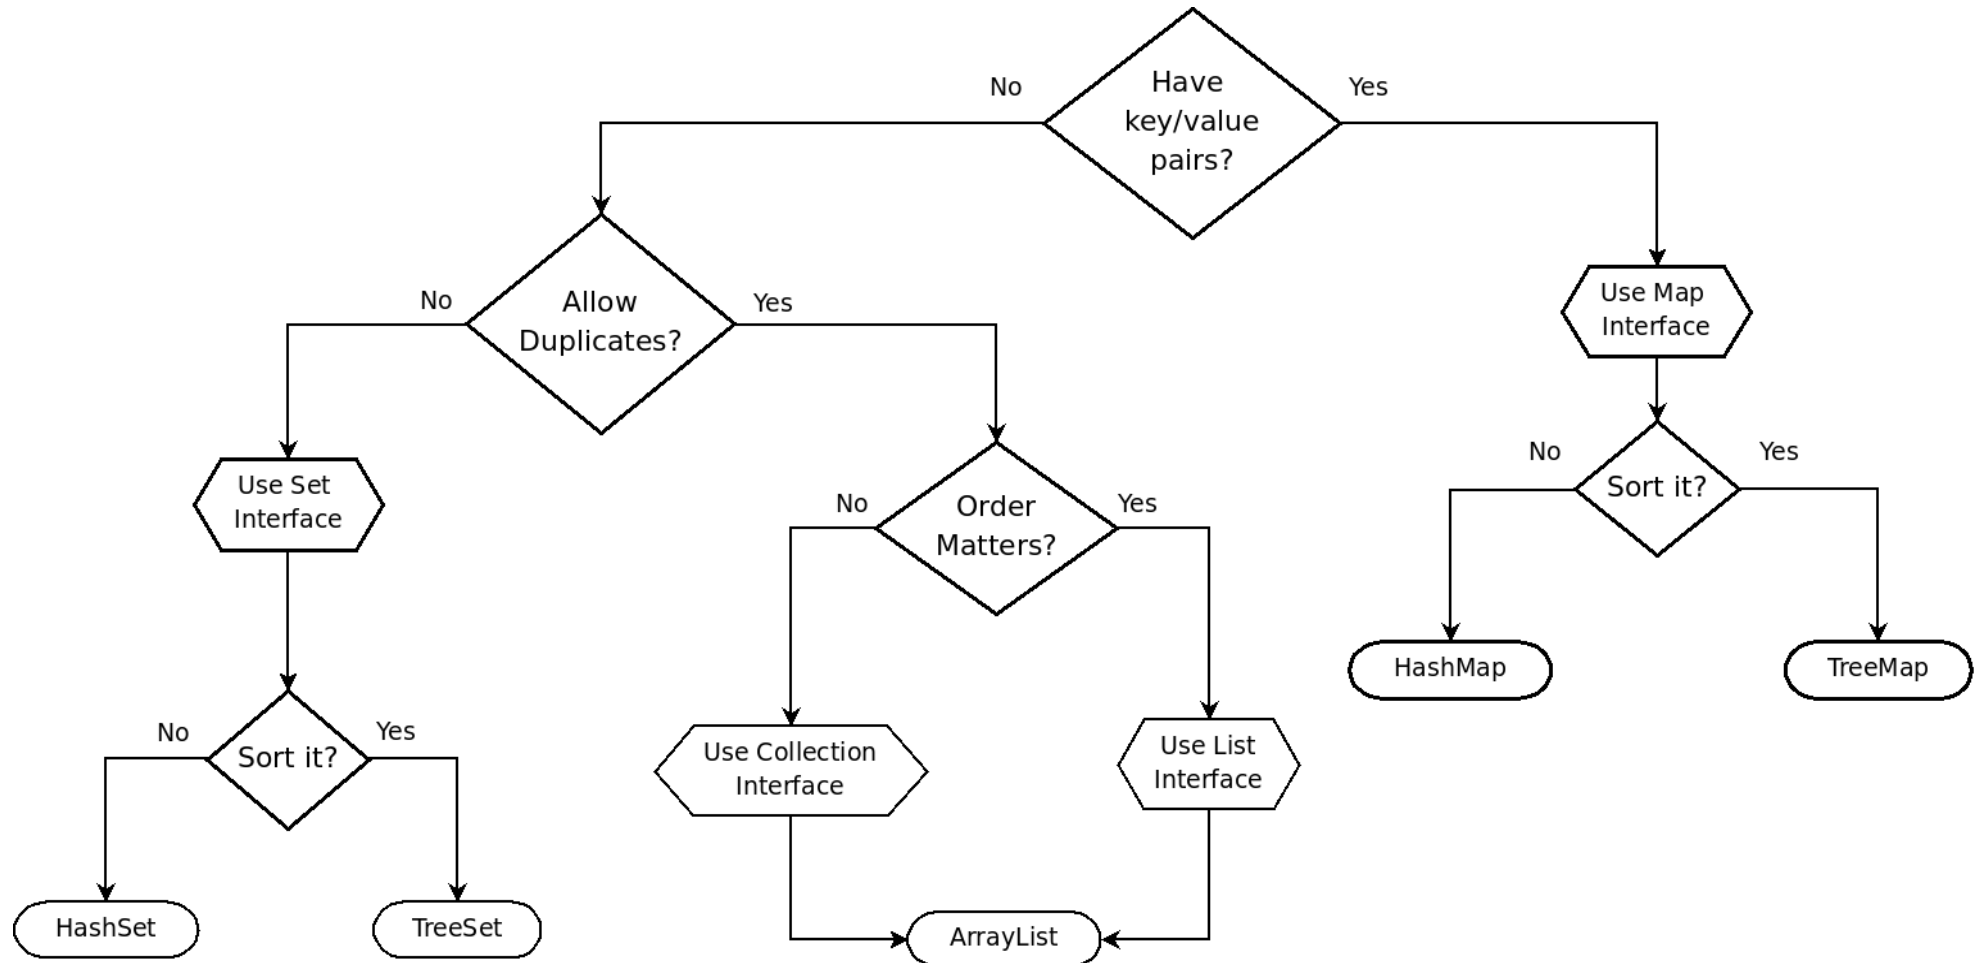
- **Collections** class (not Collection interface; note -s)
 - Provides static methods
 - First argument is a Collection to work on
 - Most apply to List, some to Collection
- Has algorithms for
 - Sorting (fast, stable version of merge sort)
 - Shuffling (randomises order of a list)
 - Routine data manipulation (reverse, fill, copy, swap, addAll)
 - Searching
 - Composition
 - frequency() counts number of times element occurs
 - disjoint() determines whether two collections are disjoint
 - Finding extreme values (max and min)

Choosing a Collection

Choosing Interfaces and Classes

- Not always easy to choose the right one
- Different approaches to choosing
- Simple approach: what properties do you want the collection to have?
 - No duplicates?
 - Keys that map to values?
 - Sorted into a particular order?
- Next slide has a flowchart
 - Covers only the main cases
 - A bit simplistic

Choosing using Collection Properties



Choosing Collections: Operations

- A more sophisticated approach
- Decide which operations you need
- Think about which classes support them
 - and how efficiently
- Two kinds of operations:
 - Those that don't change the collection (queries)
 - Those that do (mutations)
- But first we quickly review linked lists and hash tables
 - If you understand how they work, you can work out which queries they handle efficiently
 - Easier than trying to memorise which queries are efficient

Linked Lists

- A linked list is built out of *nodes*
 - The first is called the *head*, the last the *tail*
- Each node holds:
 - an element
 - a reference to the next node (except the tail)
- To add an element, Java:
 - makes a new node *n*
 - makes an existing node (if there is one) refer to *n*
 - makes *n* refer to the following node (if there is one)
- We can insert *n* anywhere in the list
 - same time cost no matter where
 - assuming we're already at the right position
- To delete an element, Java:
 - makes the node *before* it to refer to the node *after* it
 - heads and tails are special cases

Hash Tables

- Hash tables are an efficient way of implementing some collections (e.g. `HashSet`, `HashMap`)
- Basically, a hash table is an array of lists
 - These lists are called buckets
 - The user does not specify an index

```
hashTable[0] = myObj; // no
```
 - Instead the index is computed automatically for each object

```
hashTable.add(myObj); // yes
```
- Looking up an entry is efficient
 - We do not need to search the entire array (and all lists)
 - Java computes the index and searches only that list

```
hashTable.contains(myObj); // searches only 1 list
```


Queries

- “What is in the collection?”
 - Supported by all collections
 - Get an iterator and go through elements one at a time
- “Is x in the collection?”
 - Supported by all collections
 - Efficient with hash tables, less efficient with binary search of sorted collections, inefficient using iteration

Queries

- “What value is associated with x?”
 - Only supported by maps
 - Efficient (implemented with hash table)
- “How many of these are in the collection?”
 - Supported by all using iterating and counting (inefficient)
 - Special cases that are efficient:
 - Sets either have 0 or 1
 - Maps can sometimes use special value objects:
 - value is a collection storing all the relevant objects, or
 - value is an integer indicating the number of objects (instead of storing them)

Positional Queries

- “What is at integer position x?”
 - Only supported by List
 - Efficient with ArrayList, inefficient with LinkedList
- “What integer position does x have?”
 - Only supported by List
 - Requires iteration over collection (inefficient)
- “What comes before/after x?”
 - I.e. search for x, then find element before/after it
 - Supported by List, TreeMap & TreeSet
 - Use iteration to find out (inefficient)
- “What comes next?”
 - List, TreeMap & TreeSet support iteration in a defined order
 - List also supports moving iterator backwards
 - Iteration starts at the beginning
 - List can also start at the end, or at an integer position

Mutating Collections

- Insert/delete
 - ArrayList: inefficient except at end
 - But forcing it to grow is inefficient
 - LinkedList: efficient*
 - Sorted collections: not very efficient as sorting is needed
 - Efficient with collections based on hash tables
- Mutating elements is efficient*
 - Assignment to an ArrayList index points it at a new object
 - `map.entry.setValue` points key to a new value
 - Mutating internal state of an element

*Finding the position/element to work on may not be efficient!

Mutating Collections

- Swapping position of elements
 - Efficient with ArrayList (use 2 assignments)
 - Efficient with hashMap (use setValue twice)
 - Sorted collections use insert/delete
 - less efficient
 - LinkedList uses insert/delete
 - inefficient

Miscellanea

Hash Codes

- An object's index in a hash table is called its *hash code*
 - Computed by the `hashCode` method
 - In maps, only keys are hashed, not values
 - Search is faster when lists are shorter
 - In extreme case, array has length 1 and all element are in the same list
 - This is equivalent to using a list instead of a hash table!
 - So we want to distribute elements to lists more or less equally
- `hashCode` is implemented by class `Object`
 - Based on memory location of the object
 - But manipulated to fit length of array
 - So 2 different objects (usually) have different hash codes
 - Sometimes you want a different behaviour

hashCode() and equals()

- Sometimes we want different objects to be equivalent

```
Rectangle r1 = new Rectangle(0,0,5,5);  
Rectangle r2 = new Rectangle(0,0,5,5);  
r1.equals(r2); // should be true
```

- We need to override `Object's equals()` to do this
- Similarly, if we want 2 objects to have the same hash code we must override `Object's hashCode()`
 - E.g. a set should not contain both `r1` and `r2`
 - They must have same hash code, or e.g. `contains` will not search the right list
 - Compute hash code from object's state instead of address
- Check the `Object` class in the API for documentation

Rules on hashCode() and equals()

- Two objects that are equivalent must have the same hash code
 - (But not all objects with same hash code are equivalent)
- The default `equals()` and `hashCode()` use identity (memory address)
 - If you override `equals()` to use equality you must override `hashCode()` to do the same
 - If you do NOT override `equals()`, do NOT override `hashCode()`
 - I.e. make “`hashCode` consistent with `equals`”
- In a `HashMap` only keys are used to look things up
 - Values do not need a hash code
- A special case: overriding `equals` in a subclass when the superclass already overrides it
 - See <http://www.ibm.com/developerworks/java/library/j-jtp05273.html>

New Syntax in Java 7

- Java 7 is due in late 2010
- Type parameters in constructors can be inferred from variable types e.g.

```
ArrayList<String> friends = new ArrayList<>();
```

```
Map<String, Long> phoneBook = new HashMap<>();
```

- **[]** syntax for accessing ArrayList

- Allows us to write

```
friends[0] = "Rachel";
```

instead of


```
friends.set(0, "Rachel");
```

and

```
System.out.println(friends[0]);
```

instead of

```
System.out.println(friends.get(0));
```



Java 7: Collection Literals

- A literal is code that defines an immutable value
 - Unlike a variable, which can change
- Before Java 7 we could make string and number literals
- Java 7 allows collection literals of type List, Set and Map
 - Constructs and initialises collection at same time
 - These collections are immutable

```
["Rachel", "Caroline", "Sophie"]; // a List
```

```
{"Rachel", "Caroline", "Sophie"}; // a Set
```

```
{"Rachel" : 1, "Caroline" : 2, "Sophie" : 3}; // a Map
```

Using Collection Literals

- Use as a parameter to a method

```
names.addAll(["Rachel", "Caroline", "Sophie"]);
```



- As parameter to a constructor

```
ArrayList<String> newFriends = new  
    ArrayList<>(["Sam"]);
```

- Assign to a variable

- Type of variable must be the relevant interface type
- We don't actually know the implementation type!

```
List<String> friends = ["Penny", "Imi"];
```



Older Classes

- The JCF was new in Java 1.2
 - Pre-1.2 code uses older classes (“legacy collections”)
 - Some library code uses these old classes
 - They are not deprecated; there are still useable
 - But only use them if you need to
 - See the “Interoperability” part of the Collections Trail
- `Hashtable`
 - Note lowercase `t`
 - A synchronized (thread-safe) version of `HashMap`
 - Only use if you use threads
- `Vector`
 - A synchronized version of `ArrayList`
 - Only use if you use threads

More Legacy Collections

- `Enumeration`
 - Replaced by `Iterator`
 - Only use if you're working on old code that uses it
- `arrays`
 - Still useful and widely used
 - but sometimes the newer collections are better
 - Many libraries still use them even when a collection is better
 - `Collection` interface has methods to convert to array
- `Arrays` class
 - Has static methods for manipulating arrays (sorting...)
 - Like `Collections` does for collections
 - `asList` returns a list backed by the array

Conclusion

Key Points

- Adding and removing can confuse existing iterators
 - Fail-fast iterators throw `ConcurrentModificationException` *if* they detect their collection has been modified in an unauthorised way
- `ArrayList` and `LinkedList` both implement `List`
 - Each is inefficient for some uses
 - Usually you want an `ArrayList`
- `equals()` is important
 - Sets define duplicates with `equals()`
 - For sorted collections and hash tables to work `equals` must be defined correctly
- `hashCode` must be “consistent with `equals`”
 - `equals()` and `hashCode()` must both test for identity or both test for equality
- Sorting with `Comparator` and `Comparable` must be “consistent with `equals`”

Key Points

- The hash code of a key must not change
 - Otherwise the value will be lost
 - Do not mutate the state of a key if hash code is based on state
- Do not introduce duplicates into sets by mutating elements
 - It's safest to keep elements immutable
- The general-purpose JCF implementations have certain properties (allow null elements, fail-fast iterators etc.)
 - Other implementations are different – be careful!

Exercises

- Assume you have code to read a book from a file
- Think about how to implement the following:
 - List the words used in the book
 - i.e. each word only appears once in your list
 - List them in alphabetic order
 - Also list the number of times each word appears
 - Also list the pages each word appears on
- The last is called a concordance
 - People used to do them by hand!
 - This could take years

Layers of JCF Understanding

- The JCF is a minefield
 - You may have used a collection successfully in the past
 - But often doing something a little different causes a bug
 - e.g. changing state of an element can cause problems
 - The solution is often complex
- Roughly, there are 3 levels of knowledge
 - Basic: standard use of ArrayList, HashMap, Set and Iterator
 - Covered in *Objects First* and most textbooks
 - Iterators vs. for-loops with indexes: <http://en.wikipedia.org/wiki/Iterator>
 - Intermediate: more details about special cases
 - This lecture
 - The trail on which this lecture is based
 - For a JCF overview and design FAQ see also <http://java.sun.com/javase/6/docs/technotes/guides/collections/index.html>
 - *Core Java* is a good reference book with a JCF chapter

Advanced Resources

- Advanced: Full details about all the special cases, efficiency issues, implementing your own collections...
 - Specialised books that go into the JCF in detail e.g.
 - *Java Generics and Collections* by Maurice Naftalin and Philip Wadler
 - *Data Structures and the Java Collections Framework* by William Collins
 - I haven't read either...
 - Many resources on the web
 - many articles on specific subjects
 - just search for them
 - e.g. equals and hashCode
 - <http://www.ibm.com/developerworks/java/library/j-jtp05273.html>

Related Resources

- There are textbooks on data structures with example Java code
 - But they mostly show you how to implement your own data structures
 - Not much about how to use the JCF
- See *Big Java* on how to implement hash tables and linked lists yourself
 - You are not likely to need this since the JCF provides them
 - But it can be useful to understand how they work